

Introduction Our goal was to examining the relationship between stock price and google trends “search interest” across various companies. This is scaled back from our original goal to visualize the relationship between stock price and social media interest. We achieved this goal by creating a graph which simultaneously plotted the google trend data and a company’s historical stock price, then added the ability to switch which company is displayed. We decided that it would be helpful if users could see events which occurred around the same time as shifts in stock price or search interest, so we added functionality so that if one clicks on a data point, the visualization displays New York Times articles related to the company that appeared in the month leading up to the date of the data point. We also added functionality so that when one of these data points is moused over the legend will display the exact data that the data point contains. Over the course of implementing this project, we learned that the javascript date library can be difficult to work with, that good data can be hard to find, and that the relationship between stock price and public interest can often be difficult to fully understand.

Motivation As social media becomes increasingly important in the world and companies care more and more about their presence on social media, understanding the impact that popular attention has on companies becomes more crucial. By examining google trends “search interest”, we have a proxy for the public’s attention and we can use this to see how a company’s performance relates to its popularity over time.

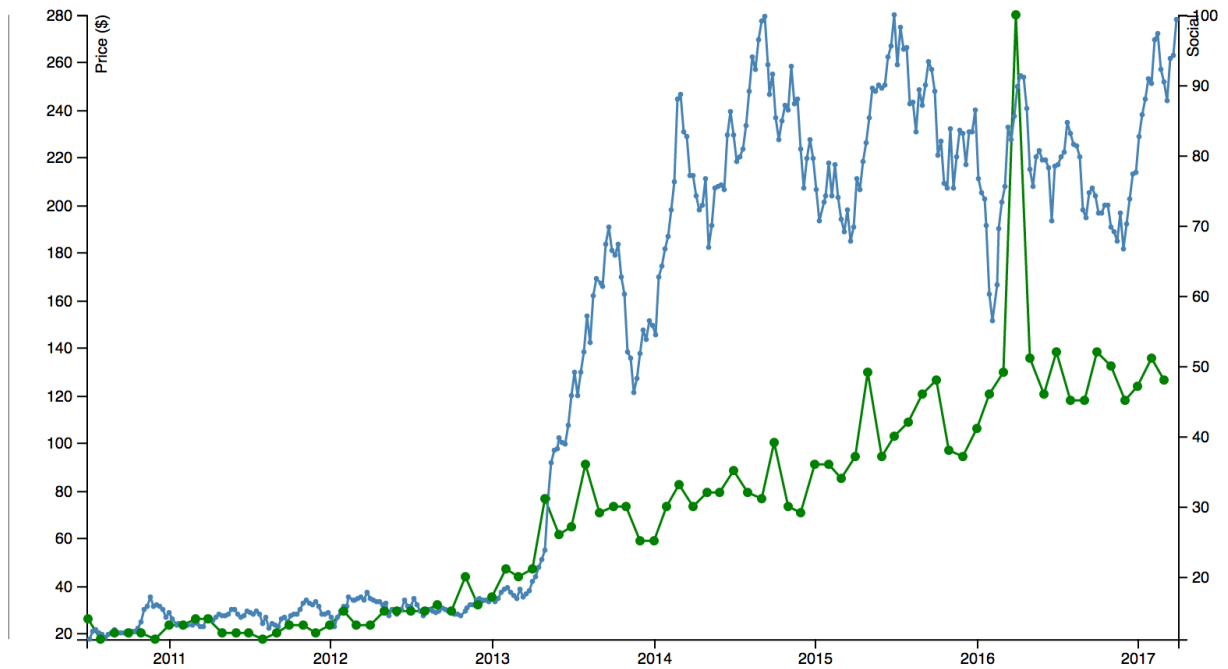
By placing this data in a visualization, we allow users to explore and compare the data available in a way which would be hard to do otherwise. It is difficult to gather this data on ones own, and without a visual comparison it is difficult to recognize how a price and an arbitrary trend scale are related. Being able to explore in a visual way is particularly important, as the relationship between stock price and social media data is not always easy to understand. Public attention can be positive or negative for a company, depending on the social and historical context. Because of this, the combination of a graph and related New York Times headlines allows users to more easily to understand the historical context of shifts in stock, or popularity. With this graph, the user can answer not only questions such as ”Is publicity generally good or bad for a company’s performance?” but also ”What events led up to such a drastic reduction in stock price and increase in social media attention for Delta Airlines in 2017?”

Although our visualization is somewhat conceptually simpler than many other visualizations, it still has the potential to answer many of our users questions. For example, if one wanted to determine whether stock prices and publicity (as measured search interest) were changing in a similar way in 2015 for Wendys as well as understand

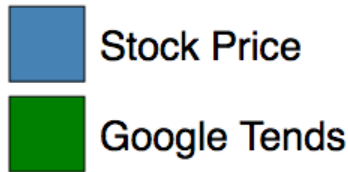
the historical context surrounding the company at that time, they could get all this information from a mere two clicks in our visualization.

Results Our visualization includes a plot with an x-axis of time and two y-axes. One y-axis is the stock values of a company, and the other is its popularity in google trends. This plot graphs changes in stock prices over time as well as changes in search popularity. An example of this for tesla is included below:

Popularity V.S. Price



One can display data for a specific company by using the dropdown menu to the left of the screen. To examine the data of a specific point, one can mouse over it. The date, and values are displayed to the right of the screen along with a legend indicating what each line indicates. Additionally, if one is curious about what events were occurring at a specific point, clicking on the point will display relevant New York Times articles from the month previous to that point in time.



Date: 7/25/2016

Closing Price: 234.78

Date: 10/1/2016

Popularity: 52

We also included an explanation tab, which contains a description of how to use that visualization, what it means, and how we got and prepared our data.

Visualization / Explanation

What this visualization does:

Our visualization includes a plot with an x-axis of time and two y-axes. One y-axis is the stock values of a company, and the other is its popularity in google trends. This plot graphs changes in stock prices over time as well as changes in search popularity.

How to use this visualization:

You can select the company for which to display data by using the menu to the left. To examine the specific data of a specific point, mouse over it. The date, and values will be displayed to the right. Additionally, if you are curious about what events were occurring at a specific point, clicking on the point will display relevant New York Times articles from around that time.

How the data was prepared:

We obtained the stock data we used for this visualization from Yahoo, and the trends data from Google Trends. When the data didn't specify when the day of the month as part of the date, we arbitrarily set the date to the first day of the month so that the data could still be displayed.

It is difficult to determine whether search interest and stock prices are correlated overall. Although they appear to shift in similar ways in many different scenarios, there are many times in which they appear unrelated. We believe that we currently have too many variables unaccounted for and too little data to make any conclusive statements about overall correlation, although we did notice that the relationship was more pronounced for some companies than for others.

Implementation Our original design was very similar to our current design, but with some additional

features that we didnt have the data to implement. We originally intended to include options for the user to display multiple lines, indicating the publicity of a company on different social media sources, and to allow the users observe the value of all of these sources, or filter out a select few of them, and to filter out a select time period to look at more closely. We also originally intended to display a map of the world, where the countries and states were highlighted based on how many of the searches came from that country. The final difference between our original and our final product is that our original design did not originally include the option of displaying new york times articles from around the time period we are looking at.

We originally intended for the user to be able to switch between different interest indicators, such as Reddit, Twitter and Facebook data, but we were unable to obtain what we were looking for from these sources. Many of them charged a large fee for their data, if it was publicly available at all, and we didnt have the monetary resources to expend on such a venture. Thus, our current product includes only data from google trends, which cost no money and was easier to obtain. We also realized soon that the New York Times had a free api that would allow us to access snippets of stories in a given time period. After becoming frustrated with the availability of Social media data, we decided to stick to using Google trends data, and added the ability to access New York times article from the relevant time.

After we had found the data we wanted to use we started to work on the implementation, and we encountered several issues that we had to work through. First, dates proved to be more difficult to work with than we intended. Creating a range of dates in javascript in order to use with the New York Times API was not well documented, and our original attempt to create a date range resulted in several bugs. In addition, we had some difficulty we didnt expect with our stock and google trends data. The Amazon stocks dates were coded differently from all the other datasets we had, using slashes instead of dashes to separate out the day month and year. Because of this were thus forced to manually change the data to a format that our code could parse. This was done relatively quickly using Emacs macros, but discovering the source of the error took some time. We also had several bugs that we eventually realized came from some google trends dates only having the year and month and not the day. We fixed this by manually adding in the date 01. Although this fixed the problem of parsing data, it also left us with slightly unsure of the accuracy of our data. Because the original data didnt include the dates, we didnt know if our data was shifted earlier, and if so by how much. We accounted this by concern by explaining what we did in the explanation part of the visualization.

The main way in which we used d3 in our visualization to make our visualization dynamic, so that users can interact with it. We feel that this makes the user feel more connected with the data, and lets them choose where and how to focus their attention. We also used d3 to add subtle transitions, which we feel make our visualization more aesthetically appealing. In the same vein, we used additional d3 libraries in order to provide structure and layout to the page, furthering its appeal. Our visualization is made almost entirely of techniques which we learned in class, although the lines we actually drew in our graph, and the vertical line which appears when hovering over the

graph is not based off our work in class. We also learned how to work with the New York Times headlines, which are based in an outside API.

Reflection With more time, there are several potential enhancements that we could have made to our visualization. First, we could have added the ability to view and compare multiple companies on the same graph. This would have given users another dimension of analysis with which to examine the data, and made it easier to search for overall correlations. Additionally, we considered using google trends to add a map of search locations, to help users understand how the data fits in spatially and whether there are any geographic trends. Lastly, we could have made minor visual fixes and enhancements, such as adding more consistent animation in the New York Times data, making sure that everything was consistent across browsers, and making adjustments so that the page scaled more appropriately.

With more data we would have been able to improve our visualization in several ways. If we had been able to access the social media data which we had originally intended to collect, we would have been able to include more lines in the graph and better depict publicity, as well as show if there were discrepancies in amount of publicity across sites. Additionally, we would have included more data points for google trends, so we wouldn't have to be concerned about the error cause by shifting the data points to the first day of the month. Finally, if we had access to the data, we would have added different metrics for evaluating the success of the companies that we looked at. Looking at metrics like quarterly evaluations, and dividends would have made it easier for the user to separate out the effects of publicity from the effects of increased profitability, as well as look at the connection between publicity and profitability for different countries.

Conclusion Our current project provides a basic but easy to use tool for exploring the relationship between publicity and stock prices. Because of restrictions on data we had access too, and the amount of time we had, we ran into limits on the sophistication of our model. However, the current visualization provides the data organized in a very comprehensible way, and allows users to easily explore the connection between publicity and stock price in the way we originally intended. Additionally, because we include New York Times articles from the relevant times, users can explore the context surrounding the data in a way beyond the original scope of our project.